

Question 1: Wavelets, Sketches**(1 P.)**

- a) Calculate the wavelet transform of the following cumulative distribution over the numbers 0 to 7. Also calculate the Mean-Squared-Error (MSE) if you omit all coefficients ≤ 3 .

[2, 4, 5, 5, 12, 20, 26, 30]

- b) Given the following Table with elements $x \in X$ and the hash values for the hash function $h : X \rightarrow [0, 1]$.

x	$h(x)$	x	$h(x)$
Toyota	0.05	Mercedes	0.14
Ford	0.47	Dodge	0.565
Bughatti	0.72	Porsche	0.63
BMW	0.26	Nissan	0.395
Audi	0.8	Hyundai	0.93

Using the KMV Sketch algorithm for $k \in 1, 2, 3, 4, 5$, calculate the estimate as well as the absolute error. Will the error get smaller with a larger k ? How does your observation explain the impact of k on KMV Sketch?

Question 2: Statistics in PostgreSQL**(1 P.)**

To improve the query optimization, PostgreSQL keeps track of various statistics about every column in every table. Consider the following statistics about the *lineitem* table from the TPC-H database:

```
SELECT * FROM pg_stats WHERE tablename = 'lineitem'
```

attname	null_frac	n_distinct	most_common_vals	most_common_freqs
l_linenumber	0	7	{1,2,3,4,5,6,7}	{0.25,0.21,0.17,0.14,0.10,0.07,0.03}
l_orderkey	0	419664	{1703939,3644579,507137,703172,712326,770882,971014}	{0.000133333,0.000133333,0.0001,0.0001,0.0001,0.0001,0.0001}
l_extendedprice	0	-0.11964	{13747.3,25165.2,26677.8,32274.0,33265.3,35938.8,50370.3 }	{0.0001,0.0001,0.0001,0.0001,0.0001,0.0001,0.0001 }

- a) What does each column of *pg_stats* represent?
- b) Which of the following queries benefit from these statistics? The total size of the table is 6 001 215.
- ☐ `SELECT * FROM lineitem`
 - ☐ `SELECT * FROM lineitem WHERE l_orderkey IS NOT NULL`
 - ☐ `SELECT * FROM lineitem WHERE l_orderkey = 406631`
 - ☐ `SELECT * FROM lineitem WHERE l_extendedprice = 32274.0`
 - ☐ `SELECT * FROM lineitem WHERE l_extendedprice <= 9500.0`
- c) Use the statistics to estimate the result size of each query in part (b).

In the table below an additional column of the *pg_stats* table is shown. It shows "A list of values that divide the column's values into groups of approximately equal population. The values in *most_common_vals*, if present, are omitted from this histogram calculation." – <https://www.postgresql.org/docs/10/static/view-pg-stats.html>

attname	histogram_bounds
l_extendedprice	{906.00,1528.60,2062.06,3035.16,3639.44,4444.04,5265.52,5920.60,6751.85,7518.63,8252.80,9032.32,9777.68,10540.68,11267.46,11921.76,12652.00,13413.40,14078.28 (...) }

- d) Decide for which queries in (b) this additional column is useful and estimate the result set size.

Question 3: Join ordering

(1 P.)

Given the following query on the TPC-H database:

```
SELECT *  
FROM orders o, customer c, nation n,  
WHERE o.o_custkey = c.c_custkey AND c.c_nationkey = n.n_nationkey
```

- Draw the query graph.
- Calculate the join selectivities.
- Calculate the cost C_{out} for all possible join trees without cross products.

Question 4: Join Ordering

(1 P.)

- a) **Left-deep trees vs. Bushy trees**

Given the following relations: R_1 , R_2 , R_3 and R_4 , with $|R_1| = 200$, $|R_2| = 100$, $|R_3| = 150$, $|R_4| = 30$, and the following join selectivities: $f_{1,2} = 0.02$, $f_{1,3} = 0.15$, $f_{1,4} = 0.08$.

- Draw the query graph.
 - List all left-deep trees without cross products.
 - Calculate the C_{out} cost for the join orderings created in (ii) where R_2 is joined last (on the top of the tree).
- b) Is there a bushy tree (that is not a linear tree), without cross products, that has lower cost?

Question 5: Join Ordering Algorithms

(1 P.)

Write a program which compares the different greedy join ordering algorithms, presented in the lecture. You may use any language you like. You have to write one method for each greedy algorithm plus two methods that find the best and worst possible left-deep join ordering. For the last two functions you have to find all possible orderings (brute force). Do not generate join orderings that contain cross-products for the best and worst methods. Use the C_{out} cost as weight function where applicable and find a suitable alternative where not. Compare the algorithm results with these costs.

You may use the Java template provided in OLAT. The template provides all required boilerplate code and instructions on how to use it.

Use your program to generate plans for the following relations: $|R_1| = 250$, $|R_2| = 150$, $|R_3| = 40$, $|R_4| = 220$ with the selectivities: $f_{1,3} = 0.01$, $f_{2,3} = 0.02$, $f_{3,4} = 0.1$

Submit the created plans with their C_{out} costs and the code that created them.

HINT: The best and worst plans (without cross products) range somewhere between (7 000, 10 500).